

Machine Learning with Python and H2O

SPENCER AIELLO CLIFF CLICK

HANK ROARK LUDI REHAK

EDITED BY: JESSICA LANFORD

<http://h2o.ai/resources/>

February 2016: Third Edition

Machine Learning with Python and H2O
by Spencer Aiello, Cliff Click,
Hank Roark & Ludi Rehak
Edited by: Jessica Lanford

Published by H2O.ai, Inc.
2307 Leghorn St.
Mountain View, CA 94043

©2016 H2O.ai, Inc. All Rights Reserved.

February 2016: Third Edition

Photos by ©H2O.ai, Inc.

All copyrights belong to their respective owners.
While every precaution has been taken in the
preparation of this book, the publisher and
authors assume no responsibility for errors or
omissions, or for damages resulting from the
use of the information contained herein.

Printed in the United States of America.

Contents

1	Introduction	4
2	What is H2O?	5
2.1	Example Code	6
2.2	Citation	6
3	Installation	6
3.1	Installation in Python	6
4	Data Preparation	7
4.1	Viewing Data	10
4.2	Selection	11
4.3	Missing Data	13
4.4	Operations	14
4.5	Merging	16
4.6	Grouping	18
4.7	Using Date and Time Data	19
4.8	Categoricals	20
4.9	Loading and Saving Data	21
5	Machine Learning	22
5.1	Modeling	22
5.1.1	Supervised Learning	23
5.1.2	Unsupervised Learning	23
5.2	Running Models	24
5.2.1	Gradient Boosting Models (GBM)	24
5.2.2	Generalized Linear Models (GLM)	27
5.2.3	K-means	31
5.2.4	Principal Components Analysis (PCA)	32
5.3	Grid Search	33
5.4	Integration with scikit-learn	34
5.4.1	Pipelines	34
5.4.2	Randomized Grid Search	36
6	References	39

1 Introduction

This documentation describes how to use H2O from Python. More information on H2O's system and algorithms (as well as complete Python user documentation) is available at the H2O website at <http://docs.h2o.ai>.

H2O Python uses a REST API to connect to H2O. To use H2O in Python or launch H2O from Python, specify the IP address and port number of the H2O instance in the Python environment. Datasets are not directly transmitted through the REST API. Instead, commands (for example, importing a dataset at specified HDFS location) are sent either through the browser or the REST API to perform the specified task.

The dataset is then assigned an identifier that is used as a reference in commands to the web server. After one prepares the dataset for modeling by defining significant data and removing insignificant data, H2O is used to create a model representing the results of the data analysis. These models are assigned IDs that are used as references in commands.

Depending on the size of your data, H2O can run on your desktop or scale using multiple nodes with Hadoop, an EC2 cluster, or Spark. Hadoop is a scalable open-source file system that uses clusters for distributed storage and dataset processing. H2O nodes run as JVM invocations on Hadoop nodes. For performance reasons, we recommend that you do not run an H2O node on the same hardware as the Hadoop NameNode.

H2O helps Python users make the leap from single machine based processing to large-scale distributed environments. Hadoop lets H2O users scale their data processing capabilities based on their current needs. Using H2O, Python, and Hadoop, you can create a complete end-to-end data analysis solution.

This document describes the four steps of data analysis with H2O:

1. installing H2O
2. preparing your data for modeling
3. creating a model using simple but powerful machine learning algorithms
4. scoring your models

2 What is H2O?

H2O is fast, scalable, open-source machine learning and deep learning for smarter applications. With H2O, enterprises like PayPal, Nielsen Catalina, Cisco, and others can use all their data without sampling to get accurate predictions faster. Advanced algorithms such as deep learning, boosting, and bagging ensembles are built-in to help application designers create smarter applications through elegant APIs. Some of our initial customers have built powerful domain-specific predictive engines for recommendations, customer churn, propensity to buy, dynamic pricing, and fraud detection for the insurance, healthcare, telecommunications, ad tech, retail, and payment systems industries.

Using in-memory compression, H2O handles billions of data rows in-memory, even with a small cluster. To make it easier for non-engineers to create complete analytic workflows, H2O's platform includes interfaces for R, Python, Scala, Java, JSON, and CoffeeScript/JavaScript, as well as a built-in web interface, Flow. H2O is designed to run in standalone mode, on Hadoop, or within a Spark Cluster, and typically deploys within minutes.

H2O includes many common machine learning algorithms, such as generalized linear modeling (linear regression, logistic regression, etc.), Naïve Bayes, principal components analysis, k-means clustering, and others. H2O also implements best-in-class algorithms at scale, such as distributed random forest, gradient boosting, and deep learning. Customers can build thousands of models and compare the results to get the best predictions.

H2O is nurturing a grassroots movement of physicists, mathematicians, and computer scientists to herald the new wave of discovery with data science by collaborating closely with academic researchers and industrial data scientists. Stanford university giants Stephen Boyd, Trevor Hastie, Rob Tibshirani advise the H2O team on building scalable machine learning algorithms. With hundreds of meetups over the past three years, H2O has become a word-of-mouth phenomenon, growing amongst the data community by a hundred-fold, and is now used by 30,000+ users and is deployed using R, Python, Hadoop, and Spark in 2000+ corporations.

Try it out

- Download H2O directly at <http://h2o.ai/download>.
- Install H2O's R package from CRAN at <https://cran.r-project.org/web/packages/h2o/>.
- Install the Python package from PyPI at <https://pypi.python.org/pypi/h2o/>.

Join the community

- To learn about our meetups, training sessions, hackathons, and product updates, visit <http://h2o.ai>.
- Visit the open source community forum at <https://groups.google.com/d/forum/h2ostream>.
- Join the chat at <https://gitter.im/h2oai/h2o-3>.

2.1 Example Code

Python code for the examples in this document is located here:

https://github.com/h2oai/h2o-3/tree/master/h2o-docs/src/booklets/v2_2015/source/python

2.2 Citation

To cite this booklet, use the following:

Aiello, S., Cliff, C., Roark, H., Rehak, L., and Lanford, J. (Feb 2016). *Machine Learning with Python and H2O*. <http://h2o.ai/resources/>.

3 Installation

H2O requires Java; if you do not already have Java installed, install it from <https://java.com/en/download/> before installing H2O.

The easiest way to directly install H2O is via a Python package.

(**Note:** The examples in this document were created with H2O version 3.9.0.3378.)

3.1 Installation in Python

To load a recent H2O package from PyPI, run:

```
1 pip install h2o
```

To download the latest stable H2O-3 build from the H2O download page:

1. Go to <http://h2o.ai/download>.
2. Choose the latest stable H2O-3 build.

3. Click the “Install in Python” tab.
4. Copy and paste the commands into your Python session.

After H2O is installed, verify the installation:

```

1 import h2o
2
3 # Start H2O on your local machine
4 h2o.init()
5
6 # Get help
7 help(h2o.estimators.glm.H2OGeneralizedLinearEstimator)
8 help(h2o.estimators.gbm.H2OGradientBoostingEstimator)
9
10 # Show a demo
11 h2o.demo("glm")
12 h2o.demo("gbm")

```

4 Data Preparation

The next sections of the booklet demonstrate the Python interface using examples, which include short snippets of code and the resulting output.

In H2O, these operations all occur distributed and in parallel and can be used on very large datasets. More information about the Python interface to H2O can be found at docs.h2o.ai.

Typically, we import and start H2O on the same machine as the running Python process:

```

1 In [1]: import h2o
2
3 In [2]: h2o.init()
4
5
6 No instance found at ip and port: localhost:54321. Trying to start local jar
7     ...
8
9 JVM stdout: /var/folders/wg/3qx1qchx1jsfjqqbzmz3stj7c0000gn/T/tmpof5ZIZ/
   h2o_hank_started_from_python.out
10 JVM stderr: /var/folders/wg/3qx1qchx1jsfjqqbzmz3stj7c0000gn/T/tmpk4uayp/
   h2o_hank_started_from_python.err
11 Using ice_root: /var/folders/wg/3qx1qchx1jsfjqqbzmz3stj7c0000gn/T/tmpKylWmt
12
13
14 Java Version: java version "1.8.0_40"
15 Java(TM) SE Runtime Environment (build 1.8.0_40-b27)
16 Java HotSpot(TM) 64-Bit Server VM (build 25.40-b25, mixed mode)

```

```

17
18
19 Starting H2O JVM and connecting: ..... Connection sucessful!
20 -----
21 H2O cluster uptime:          1 seconds 591 milliseconds
22 H2O cluster version:        3.2.0.5
23 H2O cluster name:           H2O_started_from_python
24 H2O cluster total nodes:    1
25 H2O cluster total memory:    3.56 GB
26 H2O cluster total cores:    4
27 H2O cluster allowed cores:  4
28 H2O cluster healthy:        True
29 H2O Connection ip:          127.0.0.1
30 H2O Connection port:        54321
31 -----

```

To connect to an established H2O cluster (in a multi-node Hadoop environment, for example):

```

1 In[2]: h2o.init(ip="123.45.67.89", port=54321)

```

To create an H2OFrame object from a Python tuple:

```

1 In [3]: df = h2o.H2OFrame(zip(*(1, 2, 3),
2 ...:                          ('a', 'b', 'c'),
3 ...:                          (0.1, 0.2, 0.3))))
4
5 Parse Progress: [#####] 100%
6 Uploaded py9bccf8ce-c01e-40c8-bc73-b8e7e0b17c6a into cluster with 3 rows and
  3 cols
7
8 In [4]: df
9 Out[4]: H2OFrame with 3 rows and 3 columns:
10  C1  C2  C3
11  ---  ---  ---
12  1  a   0.1
13  2  b   0.2
14  3  c   0.3

```

To create an H2OFrame object from a Python list:

```

1 In [5]: df = h2o.H2OFrame(zip(*[[1, 2, 3],
2 ...:                             ['a', 'b', 'c'],
3 ...:                             [0.1, 0.2, 0.3]]))
4
5 Parse Progress: [#####] 100%
6 Uploaded py2c9ccb17-a86e-47d7-bela-a7950b338870 into cluster with 3 rows and
  3 cols
7
8 In [6]: df
9 Out[6]: H2OFrame with 3 rows and 3 columns:
10  C1  C2  C3
11  ---  ---  ---
12  1  a   0.1
13  2  b   0.2
14  3  c   0.3

```


To create an H2OFrame object from `collections.OrderedDict` or a Python dict:

```

1 In [7]: df = h2o.H2OFrame({'A': [1, 2, 3],
2   ...:                     'B': ['a', 'b', 'c'],
3   ...:                     'C': [0.1, 0.2, 0.3]})
4
5 Parse Progress: [#####] 100%
6 Uploaded py2714e8a2-67c7-45a3-9d47-247120c5d931 into cluster with 3 rows and
   3 cols
7
8 In [8]: df
9 Out[8]: H2OFrame with 3 rows and 3 columns:
10
11   A   C   B
12 --- --- ---
13  1  0.1  a
14  2  0.2  b
15  3  0.3  c

```

To create an H2OFrame object from a Python dict and specify the column types:

```

1 In [14]: df2 = h2o.H2OFrame.from_python({'A': [1, 2, 3],
2   ....:                                  'B': ['a', 'a', 'b'],
3   ....:                                  'C': ['hello', 'all', 'world'],
4   ....:                                  'D': ['12MAR2015:11:00:00', '13
5   ....:                                  column_types=['numeric', 'enum', '
        string', 'time']})
6
7 Parse Progress: [#####] 100%
8 Uploaded pyl7ealf6d-ae83-451d-ad33-89e770061601 into cluster with 3 rows and
   4 cols
9
10 In [10]: df2
11 Out[10]: H2OFrame with 3 rows and 4 columns:
12
13   A   C   B   D
14 --- --- --
15  1  hello  a  2015-03-12 11:00:00
16  2   all  a  2015-03-13 12:00:00
17  3 world  b  2015-03-14 13:00:00

```

To display the column types:

```

1 In [11]: df2.types
2 Out[11]: {u'A': u'numeric', u'B': u'string', u'C': u'enum', u'D': u'time'}

```

4.1 Viewing Data

To display the top and bottom of an H2OFrame:

```

1 In [16]: import numpy as np
2
3 In [17]: df = h2o.H2OFrame.from_python(np.random.randn(4,100).tolist(),
4     column_names=list('ABCD'))
5
6 Parse Progress: [#####] 100%
7 Uploaded py0a4d1d8d-7d04-438a-a97f-a9521f802366 into cluster with 100 rows
8 and 4 cols
9
10 In [18]: df.head()
11 H2OFrame with 100 rows and 4 columns:
12
13      A      B      C      D
14 -----
15 -0.613035 -0.425327 -1.92774 -2.1201
16 -1.26552 -0.241526 -0.0445104 1.90628
17 0.763851 0.0391609 -0.500049 0.355561
18 -1.24842 0.912686 -0.61146 1.94607
19 2.1058 -1.83995 0.453875 -1.69911
20 1.7635 0.573736 -0.309663 -1.51131
21 -0.781973 0.051883 -0.403075 0.569406
22 1.40085 1.91999 0.514212 -1.47146
23 -0.746025 -0.632182 1.27455 -1.35006
24 -1.12065 0.374212 0.232229 -0.602646
25
26 In [19]: df.tail(5)
27 H2OFrame with 100 rows and 4 columns:
28
29      A      B      C      D
30 -----
31 1.00098 -1.43183 -0.322068 0.374401
32 1.16553 -1.23383 -1.71742 1.01035
33 -1.62351 -1.13907 2.1242 -0.275453
34 -0.479005 -0.0048988 0.224583 0.219037
35 -0.74103 1.13485 0.732951 1.70306

```

To display the column names:

```

1 In [20]: df.columns
2 Out[20]: [u'A', u'B', u'C', u'D']

```

To display compression information, distribution (in multi-machine clusters), and summary statistics of your data:

```

1 In [21]: df.describe()
2 Rows: 100 Cols: 4
3
4 Chunk compression summary:
5 chunk_type      chunkname      count      count_%      size      size_%
6 -----
7 64-bit Reals      C8D      4      100      3.4 KB      100
8
9 Frame distribution summary:
10      size      #_rows      #_chunks_per_col      #_chunks
11 -----
12 127.0.0.1:54321      3.4 KB      100      1      4
13 mean      3.4 KB      100      1      4

```

```

14 min          3.4 KB  100      1          4
15 max          3.4 KB  100      1          4
16 stddev       0 B     0        0          0
17 total        3.4 KB  100      1          4
18
19 Column-by-Column Summary: (floats truncated)
20
21 -----
22      A          B          C          D
23 type    real    real    real    real
24 mins   -2.49822 -2.37446 -2.45977 -3.48247
25 maxs    2.59380  1.91998  3.13014  2.39057
26 mean   -0.01062 -0.23159  0.11423 -0.16228
27 sigma    1.04354  0.90576  0.96133  1.02608
28 zero_count  0        0        0        0
29 missing_count 0        0        0        0

```

4.2 Selection

To select a single column by name, resulting in an H2OFrame:

```

1 In [23]: df['A']
2 Out[23]: H2OFrame with 100 rows and 1 columns:
3      A
4 0 -0.613035
5 1 -1.265520
6 2  0.763851
7 3 -1.248425
8 4  2.105805
9 5  1.763502
10 6 -0.781973
11 7  1.400853
12 8 -0.746025
13 9 -1.120648

```

To select a single column by index, resulting in an H2OFrame:

```

1 In [24]: df[1]
2 Out[24]: H2OFrame with 100 rows and 1 columns:
3      B
4 0 -0.425327
5 1 -0.241526
6 2  0.039161
7 3  0.912686
8 4 -1.839950
9 5  0.573736
10 6  0.051883
11 7  1.919987
12 8 -0.632182
13 9  0.374212

```

To select multiple columns by name, resulting in an H2OFrame:

```

1 In [25]: df[['B','C']]
2 Out[25]: H2OFrame with 100 rows and 2 columns:
3           B           C
4 0 -0.425327 -1.927737
5 1 -0.241526 -0.044510
6 2  0.039161 -0.500049
7 3  0.912686 -0.611460
8 4 -1.839950  0.453875
9 5  0.573736 -0.309663
10 6  0.051883 -0.403075
11 7  1.919987  0.514212
12 8 -0.632182  1.274552
13 9  0.374212  0.232229

```

To select multiple columns by index, resulting in an H2OFrame:

```

1 In [26]: df[0:2]
2 Out[26]: H2OFrame with 100 rows and 2 columns:
3           A           B
4 0 -0.613035 -0.425327
5 1 -1.265520 -0.241526
6 2  0.763851  0.039161
7 3 -1.248425  0.912686
8 4  2.105805 -1.839950
9 5  1.763502  0.573736
10 6 -0.781973  0.051883
11 7  1.400853  1.919987
12 8 -0.746025 -0.632182
13 9 -1.120648  0.374212

```

To select multiple rows by slicing, resulting in an H2OFrame:

Note By default, H2OFrame selection is for columns, so to slice by rows and get all columns, be explicit about selecting all columns:

```

1 In [27]: df[2:7, :]
2 Out[27]: H2OFrame with 5 rows and 4 columns:
3           A           B           C           D
4 0  0.763851  0.039161 -0.500049  0.355561
5 1 -1.248425  0.912686 -0.611460  1.946068
6 2  2.105805 -1.839950  0.453875 -1.699112
7 3  1.763502  0.573736 -0.309663 -1.511314
8 4 -0.781973  0.051883 -0.403075  0.569406

```

To select rows based on specific criteria, use Boolean masking:

```

1 In [28]: df2[ df2["B"] == "a", :]
2 Out[28]: H2OFrame with 2 rows and 4 columns:
3           A           C           B           D
4 0 1  hello  a 2015-03-12 11:00:00
5 1 2   all  a 2015-03-13 12:00:00

```

4.3 Missing Data

The H2O parser can handle many different representations of missing data types, including '' (blank), 'NA', and None (Python). They are all displayed as NaN in Python.

To create an H2OFrame from Python with missing elements:

```

1 In [46]: df3 = h2o.H2OFrame.from_python(
2     {'A': [1, 2, 3, None, ''],
3       'B': ['a', 'a', 'b', 'NA', 'NA'],
4       'C': ['hello', 'all', 'world', None, None],
5       'D': ['12MAR2015:11:00:00', None,
6            '13MAR2015:12:00:00', None,
7            '14MAR2015:13:00:00']},
8     column_types=['numeric', 'enum', 'string', 'time'])
9
10 In [47]: df3
11 Out[47]: H2OFrame with 5 rows and 4 columns:
12   A      C      B      D
13 0  1  hello    a  1.426183e+12
14 1  2   all    a           NaN
15 2  3 world    b  1.426273e+12
16 3 NaN   NaN   NaN           NaN
17 4 NaN   NaN   NaN  1.426363e+12

```

To determine which rows are missing data for a given column ('1' indicates missing):

```

1 In [49]: df3["A"].isna()
2 Out[49]: H2OFrame with 5 rows and 1 columns:
3   C1
4 0  0
5 1  0
6 2  0
7 3  1
8 4  1

```

To change all missing values in a column to a different value:

```

1 In [52]: df3
2 Out[52]: H2OFrame with 5 rows and 4 columns:
3   A      C      B      D
4 0  1  hello    a  1.426183e+12
5 1  2   all    a           NaN
6 2  3 world    b  1.426273e+12
7 3  5   NaN   NaN           NaN
8 4  5   NaN   NaN  1.426363e+12

```

To determine the locations of all missing data in an H2OFrame:

```

1 In [53]: df3.isna()
2 Out[53]: H2OFrame with 5 rows and 4 columns:
3      C1  C2  C3  C4
4  0    0    0    0    0
5  1    0    0    0    1
6  2    0    0    0    0
7  3    0    1    0    1
8  4    0    1    0    0

```

4.4 Operations

When performing a descriptive statistic on an entire H2OFrame, missing data is generally excluded and the operation is only performed on the columns of the appropriate data type:

```

1 In [60]: df3 = h2o.H2OFrame.from_python(
2     {'A': [1, 2, 3, None, ''],
3      'B': ['a', 'a', 'b', 'NA', 'NA'],
4      'C': ['hello', 'all', 'world', None, None],
5      'D': ['12MAR2015:11:00:00', None,
6           '13MAR2015:12:00:00', None,
7           '14MAR2015:13:00:00']},
8     column_types=['numeric', 'enum', 'string', 'time'])
9
10 In [61]: df4.mean(na_rm=True)
11 Out[61]: [2.0, u'NaN', u'NaN', u'NaN']

```

When performing a descriptive statistic on a single column of an H2OFrame, missing data is generally *not* excluded:

```

1 In [62]: df4["A"].mean()
2 Out[62]: [u'NaN']
3
4 In [64]: df4["A"].mean(na_rm=True)
5 Out[64]: [2.0]

```

In both examples, a native Python object is returned (list and float respectively in these examples).

When applying functions to each column of the data, an H2OFrame containing the means of each column is returned:

```

1 In [5]: df5 = h2o.H2OFrame.from_python(
2     np.random.randn(4,100).tolist(),
3     column_names=list('ABCD'))
4 Parse Progress: [#####] 100%
5
6 In [6]: df5.apply(lambda x: x.mean(na_rm=True))
7 Out[6]: H2OFrame with 1 rows and 4 columns:
8      A      B      C      D
9  0  0.020849 -0.052978 -0.037272 -0.01664

```

When applying functions to each row of the data, an H2OFrame containing the sum of all columns is returned:

```

1 In [26]: df5.apply(lambda row: sum(row), axis=1)
2 Out[26]: H2OFrame with 100 rows and 1 columns:
3          C1
4 0  0.906854
5 1  0.790760
6 2 -0.217604
7 3 -0.978141
8 4  2.180175
9 5 -2.420732
10 6  0.875716
11 7 -1.077747
12 8  2.321706
13 9 -0.700436

```

H2O provides many methods for histogramming and discretizing data. Here is an example using the `hist` method on a single data frame:

```

1 In [49]: df6 = h2o.H2OFrame(
2           np.random.randint(0, 7, size=100).tolist())
3
4 Parse Progress: [#####] 100%
5 Uploaded py5b584604-73ff-4037-9618-c53122cd0343 into cluster with 100 rows
  and 1 cols
6
7 In [50]: df6.hist(plot=False)
8
9 Parse Progress: [#####] 100%
10 Uploaded py8a993d29-e354-44cf-b10e-d97aa6fdfd74 into cluster with 8 rows and
   1 cols
11 Out[50]: H2OFrame with 8 rows and 5 columns:
12    breaks  counts  mids_true  mids  density
13 0    0.75     NaN         NaN    NaN  0.000000
14 1    1.50      10         0.0  1.125  0.116667
15 2    2.25       6         0.5  1.875  0.070000
16 3    3.00      17         1.0  2.625  0.198333
17 4    3.75       0         0.0  3.375  0.000000
18 5    4.50      16         1.5  4.125  0.186667
19 6    5.25      19         2.0  4.875  0.221667

```

H2O includes a set of string processing methods in the H2OFrame class that make it easy to operate on each element in an H2OFrame.

To determine the number of times a string is contained in each element:

```

1 In [62]: df7 = h2o.H2OFrame.from_python(
2           ['Hello', 'World', 'Welcome', 'To', 'H2O', 'World'])
3
4 In [63]: df7
5 Out[63]: H2OFrame with 6 rows and 1 columns:
6          C1
7 0    Hello
8 1    World
9 2  Welcome
10 3      To
11 4     H2O
12 5    World

```

```

13
14 In [65]: df7.countmatches('l')
15 Out[65]: H2OFrame with 6 rows and 1 columns:
16 C1
17 0 2
18 1 1
19 2 1
20 3 0
21 4 0
22 5 1

```

To replace the first occurrence of 'l' (lower case letter) with 'x' and return a new H2OFrame:

```

1 In [89]: df7.sub('l','x')
2 Out[89]: H2OFrame with 6 rows and 1 columns:
3 C1
4 0 Hexlo
5 1 Worxd
6 2 Wexcome
7 3 To
8 4 H2O
9 5 Worxd

```

For global substitution, use `gsub`. Both `sub` and `gsub` support regular expressions. To split strings based on a regular expression:

```

1 In [86]: df7.strsplit('(l)+')
2 Out[86]: H2OFrame with 6 rows and 2 columns:
3 C1 C2
4 0 He o
5 1 Wor d
6 2 We come
7 3 To NaN
8 4 H2O NaN
9 5 Wor d

```

4.5 Merging

To combine two H2OFrames together by appending one as rows and return a new H2OFrame:

```

1 In [98]: df8 = h2o.H2OFrame.from_python(np.random.randn(100,4).tolist(),
2      column_names=list('ABCD'))
3 Parse Progress: [#####] 100%
4 Uploaded py9607f2cc-087a-4d99-ba9f-917ca852clf2 into cluster with 100 rows
   and 4 cols
5
6 In [99]: df9 = h2o.H2OFrame.from_python(
7      np.random.randn(100,4).tolist(),
8      column_names=list('ABCD'))
9
10 Parse Progress: [#####] 100%
11 Uploaded pycb8b3aba-77d6-4383-88dd-4729f1f2c314 into cluster with 100 rows
   and 4 cols

```



```

12
13 In [100]: df8.rbind(df9)
14 Out[100]: H2OFrame with 200 rows and 4 columns:
15      A      B      C      D
16 0 -0.095807  0.944757  0.160959  0.271681
17 1 -0.950010  0.669040  0.664983  1.535805
18 2  0.172176  0.657167  0.970337 -0.419208
19 3  0.589829 -0.516749 -1.598524 -1.346773
20 4  1.044948 -0.281243 -0.411052  0.959717
21 5  0.498329  0.170340  0.124479 -0.170742
22 6  1.422841 -0.409794 -0.525356  2.155962
23 7  0.944803  1.192007 -1.075689  0.017082

```

For successful row binding, the column names and column types between the two H2OFrames must match.

H2O also supports merging two frames together by matching column names:

```

1 In [108]: df10 = h2o.H2OFrame.from_python( {
2           'A': ['Hello', 'World',
3               'Welcome', 'To',
4               'H2O', 'World'],
5           'n': [0,1,2,3,4,5]} )
6
7 Parse Progress: [#####] 100%
8 Uploaded py57e84cb6-ce29-4d13-afe4-4333b2186c72 into cluster with 6 rows and
9   2 cols
10
11 In [109]: df11 = h2o.H2OFrame.from_python(np.random.randint(0, 10, size=100).
12   tolist9, column_names=['n'])
13
14 Parse Progress: [#####] 100%
15 Uploaded py090fa929-b434-43c0-81bd-b9c61b553a31 into cluster with 100 rows
16   and 1 cols
17
18 In [112]: df11.merge(df10)
19 Out[112]: H2OFrame with 100 rows and 2 columns:
20      n      A
21 0 7    NaN
22 1 3     To
23 2 0  Hello
24 3 9    NaN
25 4 9    NaN
26 5 3     To
27 6 4    H2O
28 7 4    H2O
29 8 5  World
30 9 4    H2O

```

4.6 Grouping

"Grouping" refers to the following process:

- splitting the data into groups based on some criteria
- applying a function to each group independently
- combining the results into an H2OFrame

To group and then apply a function to the results:

```

1 In [123]: df12 = h2o.H2OFrame(
2   {'A' : ['foo', 'bar', 'foo', 'bar',
3         'foo', 'bar', 'foo', 'foo'],
4     'B' : ['one', 'one', 'two', 'three',
5         'two', 'two', 'one', 'three'],
6     'C' : np.random.randn(8),
7     'D' : np.random.randn(8)})
8
9 Parse Progress: [#####] 100%
10 Uploaded pyd297bab5-4e4e-4a89-9b85-f8fecf37f264 into cluster with 8 rows and
   4 cols
11
12 In [124]: df12
13 Out[124]: H2OFrame with 8 rows and 4 columns:
14
15   A      C      B      D
16 0 foo  1.583908  one -0.441779
17 1 bar  1.055763  one  1.733467
18 2 foo -1.200572  two  0.970428
19 3 bar -1.066722 three -0.311055
20 4 foo -0.023385  two  0.077905
21 5 bar  0.758202  two  0.521504
22 6 foo  0.098259  one -1.391587
23 7 foo  0.412450 three -0.050374
24
25 In [125]: df12.group_by('A').sum().frame
26 Out[125]: H2OFrame with 2 rows and 4 columns:
27
28   A      sum_C  sum_B  sum_D
29 0 bar  0.747244     3  1.943915
30 1 foo  0.870661     5 -0.835406

```

To group by multiple columns and then apply a function:

```

1 In [127]: df13 = df12.group_by(['A', 'B']).sum().frame
2
3 In [128]: df13
4 Out[128]: H2OFrame with 6 rows and 4 columns:
5
6   A      B      sum_C  sum_D
7 0 bar  one  1.055763  1.733467
8 1 bar  two  0.758202  0.521504
9 2 foo three 0.412450 -0.050374
10 3 foo  one  1.682168 -1.833366
11 4 foo  two -1.223957  1.048333
12 5 bar three -1.066722 -0.311055

```

To join the results into the original H2OFrame:

```

1 In [129]: df12.merge(df13)
2 Out[129]: H2OFrame with 8 rows and 6 columns:
3      A      B      C      D      sum_C      sum_D
4 0  foo   one  1.583908 -0.441779  1.682168 -1.833366
5 1  bar   one  1.055763  1.733467  1.055763  1.733467
6 2  foo   two -1.200572  0.970428 -1.223957  1.048333
7 3  bar  three -1.066722 -0.311055 -1.066722 -0.311055
8 4  foo   two -0.023385  0.077905 -1.223957  1.048333
9 5  bar   two  0.758202  0.521504  0.758202  0.521504
10 6  foo   one  0.098259 -1.391587  1.682168 -1.833366
11 7  foo  three  0.412450 -0.050374  0.412450 -0.050374

```

4.7 Using Date and Time Data

H2O has powerful features for ingesting and feature engineering using time data. Internally, H2O stores time information as an integer of the number of milliseconds since the epoch.

To ingest time data natively, use one of the supported time input formats:

```

1 In [140]: df14 = h2o.H2OFrame.from_python(
2           {'D': ['18OCT2015:11:00:00',
3                 '19OCT2015:12:00:00',
4                 '20OCT2015:13:00:00']},
5           column_types=['time'])
6
7 In [141]: df14.types
8 Out[141]: {'u'D': u'time'}

```

To display the day of the month:

```

1 In [142]: df14['D'].day()
2 Out[142]: H2OFrame with 3 rows and 1 columns:
3      D
4 0  18
5 1  19
6 2  20

```

To display the day of the week:

```

1 In [143]: df14['D'].dayOfWeek()
2 Out[143]: H2OFrame with 3 rows and 1 columns:
3      D
4 0  Sun
5 1  Mon
6 2  Tue

```

4.8 Categoricals

H2O handles categorical (also known as enumerated or factor) values in an H2OFrame. This is significant because categorical columns have specific treatments in each of the machine learning algorithms.

Using 'df12' from above, H2O imports columns A and B as categorical/enumerated/factor types:

```
1 In [145]: df12.types
2 Out[145]: {u'A': u'Enum', u'B': u'Enum',
3           u'C': u'Numeric', u'D': u'Numeric'}
```

To determine if any column is a categorical/enumerated/factor type:

```
1 In [148]: df12.anyfactor()
2 Out[148]: True
```

To view the categorical levels in a single column:

```
1 In [149]: df12["A"].levels()
2 Out[149]: ['bar', 'foo']
```

To create categorical interaction features:

```
1 In [163]: df12.interaction(['A','B'], pairwise=False, max_factors=3,
2           min_occurrence=1)
3 Interactions Progress: [#####] 100%
4 Out[163]: H2OFrame with 8 rows and 1 columns:
5     A_B
6 0  foo_one
7 1  bar_one
8 2  foo_two
9 3   other
10 4  foo_two
11 5   other
12 6  foo_one
13 7   other
```

To retain the most common categories and set the remaining categories to a common 'Other' category and create an interaction of a categorical column with itself:

```

1 In [168]: bb_df = df12.interaction(['B','B'], pairwise=False, max_factors=2,
2         min_occurrence=1)
3
4 Interactions Progress: [#####] 100%
5
6 In [169]: bb_df
7 Out[169]: H2OFrame with 8 rows and 1 columns:
8     B_B
9 0    one
10 1    one
11 2    two
12 3 other
13 4    two
14 5    two
15 6    one
16 7 other

```

These can then be added as a new column on the original dataframe:

```

1 In [170]: df15 = df12.cbind(bb_df)
2
3 In [171]: df15
4 Out[171]: H2OFrame with 8 rows and 5 columns:
5     A      B      C      D      B_B
6 0 foo    one  1.583908 -0.441779    one
7 1 bar    one  1.055763  1.733467    one
8 2 foo    two -1.200572  0.970428    two
9 3 bar  three -1.066722 -0.311055  other
10 4 foo    two -0.023385  0.077905    two
11 5 bar    two  0.758202  0.521504    two
12 6 foo    one  0.098259 -1.391587    one
13 7 foo  three  0.412450 -0.050374  other

```

4.9 Loading and Saving Data

In addition to loading data from Python objects, H2O can load data directly from:

- disk
- network file systems (NFS, S3)
- distributed file systems (HDFS)
- HTTP addresses

H2O currently supports the following file types:

- CSV (delimited) files
- ARFF
- ORC
- XLS
- SVMLite
- XLST

To load data from the same machine running H2O:

```
1 In[172]: df = h2o.upload_file("/pathToFile/fileName")
```

To load data from the machine running Python to the machine running H2O:

```
1 In[173]: df = h2o.import_file("/pathToFile/fileName")
```

To save an H2OFrame on the machine running H2O:

```
1 In[174]: h2o.export_file(df, "/pathToFile/fileName")
```

To save an H2OFrame on the machine running Python:

```
1 In[175]: h2o.download_csv(df, "/pathToFile/fileName")
```

5 Machine Learning

The following sections describe some common model types and features.

5.1 Modeling

The following section describes the features and functions of some common models available in H2O. For more information about running these models in Python using H2O, refer to the documentation on the H2O.ai website or to the booklets on specific models.

H2O supports the following models:

- Deep Learning
- Generalized Linear Models (GLM)
- Naïve Bayes
- Gradient Boosted Regression (GBM)
- Principal Components Analysis (PCA)
- Distributed Random Forest (DRF)
- K-means

The list is growing quickly, so check www.h2o.ai to see the latest additions.

5.1.1 Supervised Learning

Generalized Linear Models (GLM): Provides flexible generalization of ordinary linear regression for response variables with error distribution models other than a Gaussian (normal) distribution. GLM unifies various other statistical models, including Poisson, linear, logistic, and others when using ℓ_1 and ℓ_2 regularization.

Distributed Random Forest: Averages multiple decision trees, each created on different random samples of rows and columns. It is easy to use, non-linear, and provides feedback on the importance of each predictor in the model, making it one of the most robust algorithms for noisy data.

Gradient Boosting (GBM): Produces a prediction model in the form of an ensemble of weak prediction models. It builds the model in a stage-wise fashion and is generalized by allowing an arbitrary differentiable loss function. It is one of the most powerful methods available today.

Deep Learning: Models high-level abstractions in data by using non-linear transformations in a layer-by-layer method. Deep learning is an example of supervised learning, which can use unlabeled data that other algorithms cannot.

Naïve Bayes: Generates a probabilistic classifier that assumes the value of a particular feature is unrelated to the presence or absence of any other feature, given the class variable. It is often used in text categorization.

5.1.2 Unsupervised Learning

K-Means: Reveals groups or clusters of data points for segmentation. It clusters observations into k -number of points with the nearest mean.

Principal Component Analytiss (PCA): The algorithm is carried out on a set of possibly collinear features and performs a transformation to produce a new set of uncorrelated features.

Anomaly Detection: Identifies the outliers in your data by invoking the deep learning autoencoder, a powerful pattern recognition model.

5.2 Running Models

This section describes how to run the following model types:

- Gradient Boosted Models (GBM)
- Generalized Linear Models (GLM)
- K-means
- Principal Components Analysis (PCA)

as well as how to generate predictions.

5.2.1 Gradient Boosting Models (GBM)

To generate gradient boosting models for creating forward-learning ensembles, use `H2OGradientBoostingEstimator`.

The construction of the estimator defines the parameters of the estimator and the call to `H2OGradientBoostingEstimator.train` trains the estimator on the specified data. This pattern is common for each of the H2O algorithms.

```

1 In [1]: import h2o
2
3 In [2]: h2o.init()
4
5 Java Version: java version "1.8.0_40"
6 Java(TM) SE Runtime Environment (build 1.8.0_40-b27)
7 Java HotSpot(TM) 64-Bit Server VM (build 25.40-b25, mixed mode)
8
9
10 Starting H2O JVM and connecting: ..... Connection successful!
11 -----
12 H2O cluster uptime:      1 seconds 738 milliseconds
13 H2O cluster version:    3.5.0.3238
14 H2O cluster name:       H2O_started_from_python
15 H2O cluster total nodes: 1
16 H2O cluster total memory: 3.56 GB
17 H2O cluster total cores: 4
18 H2O cluster allowed cores: 4
19 H2O cluster healthy:    True
20 H2O Connection ip:      127.0.0.1
21 H2O Connection port:    54321
22 -----
23
24 In [3]: from h2o.estimators.gbm import H2OGradientBoostingEstimator
25
26 In [4]: iris_data_path = h2o.system_file("iris.csv") # load demonstration
           data
27
28 In [5]: iris_df = h2o.import_file(path=iris_data_path)
29
30 Parse Progress: [#####] 100%
31 Imported /Users/hank/PythonEnvs/h2obleeding/bin/./h2o_data/iris.csv. Parsed
    150 rows and 5 cols

```



```

32
33 In [6]: iris_df.describe()
34 Rows:150 Cols:5
35
36 Chunk compression summary:
37 chunktype chunkname count count_% size size_%
38 -----
39 1-Byte Int C1 1 20 218B 18.890
40 1-Byte Flt C2 4 80 936B 81.109
41
42 Frame distribution summary:
43 size rows chunks/col chunks
44 -----
45 127.0.0.1:54321 1.1KB 150 1 5
46 mean 1.1KB 150 1 5
47 min 1.1KB 150 1 5
48 max 1.1KB 150 1 5
49 stddev 0 B 0 0
50 total 1.1 KB 150 1 5
51
52 In [7]: gbm_regressor = H2OGradientBoostingEstimator(distribution="gaussian",
53 ntrees=10, max_depth=3, min_rows=2, learn_rate="0.2")
54
55 In [8]: gbm_regressor.train(x=range(1,iris_df.ncol), y=0, training_frame=
56 iris_df)
57
58 gbm Model Build Progress: [#####] 100%
59
60 In [9]: gbm_regressor
61 Out[9]: Model Details
62 =====
63 H2OGradientBoostingEstimator: Gradient Boosting Machine
64 Model Key: GBM_model_python_1446220160417_2
65
66 Model Summary:
67 number_of_trees | 10
68 model_size_in_bytes | 1535
69 min_depth | 3
70 max_depth | 3
71 mean_depth | 3
72 min_leaves | 7
73 max_leaves | 8
74 mean_leaves | 7.8
75
76 ModelMetricsRegression: gbm
77 ** Reported on train data. **
78
79 MSE: 0.0706936802293
80
81 R^2: 0.896209989184
82 Mean Residual Deviance: 0.0706936802293
83
84 Scoring History:
85 timestamp duration number_of_trees training_MSE
86 -----
87 2015-10-30 08:50:00 0.121 sec 1 0.472445
88 0.472445
89 2015-10-30 08:50:00 0.151 sec 2 0.334868
90 0.334868
91 2015-10-30 08:50:00 0.162 sec 3 0.242847
92 0.242847

```

87	2015-10-30 08:50:00	0.175 sec	4	0.184128
	0.184128			
88	2015-10-30 08:50:00	0.187 sec	5	0.14365
	0.14365			
89	2015-10-30 08:50:00	0.197 sec	6	0.116814
	0.116814			
90	2015-10-30 08:50:00	0.208 sec	7	0.0992098
	0.0992098			
91	2015-10-30 08:50:00	0.219 sec	8	0.0864125
	0.0864125			
92	2015-10-30 08:50:00	0.229 sec	9	0.077629
	0.077629			
93	2015-10-30 08:50:00	0.238 sec	10	0.0706937
	0.0706937			
94	Variable Importances:			
95	variable	relative_importance	scaled_importance	percentage
96	-----	-----	-----	-----
97				
98	C3	227.562	1	0.894699
99	C2	15.1912	0.0667563	0.0597268
100	C5	9.50362	0.0417627	0.037365
101	C4	2.08799	0.00917544	0.00820926

To generate a classification model that uses labels,
use `distribution="multinomial"`:

```
1 In [10]: gbm_classifier = H2OGradientBoostingEstimator(distribution="
2           multinomial", ntrees=10, max_depth=3, min_rows=2, learn_rate="0.2")
3 In [11]: gbm_classifier.train(x=range(0,iris_df.ncol-1), y=iris_df.ncol-1,
4           training_frame=iris_df)
5 gbm Model Build Progress: [#
6           #####] 100%
7 In [12]: gbm_classifier
8 Out[12]: Model Details
9 =====
10 H2OGradientBoostingEstimator : Gradient Boosting Machine
11 Model Key: GBM_model_python_1446220160417_4
12
13 Model Summary:
14   number_of_trees  model_size_in_bytes  min_depth  max_depth
15   mean_depth      min_leaves    max_leaves  mean_leaves
16   -----
17   30              3933            1           3
18   2.93333        2              8          5.86667
19
20 ModelMetricsMultinomial: gbm
21 ** Reported on train data. **
22 MSE: 0.00976685294679
23 R^2: 0.98534972058
24 LogLoss: 0.0782480971236
25
26 Confusion Matrix: vertical: actual; across: predicted
27
28 Iris-setosa    Iris-versicolor  Iris-virginica  Error    Rate
29 -----
```

```

30 50          0          0          0          0 / 50
31 0           49         1          0.02        1 / 50
32 0           0          50         0           0 / 50
33 50          49         51         0.00666667  1 / 150
34
35 Top-3 Hit Ratios:
36 k      hit_ratio
37 ---  -----
38 1      0.993333
39 2      1
40 3      1
41
42 Scoring History:
43 timestamp      duration      number_of_trees      training_MSE
44 training_logloss      training_classification_error
45 -----
46 2015-10-30 08:51:52  0.047 sec  1      0.282326
47      0.758411      0.0266667
48 2015-10-30 08:51:52  0.068 sec  2      0.179214
49      0.550506      0.0266667
50 2015-10-30 08:51:52  0.086 sec  3      0.114954
51      0.412173      0.0266667
52 2015-10-30 08:51:52  0.100 sec  4      0.0744726
53      0.313539      0.02
54 2015-10-30 08:51:52  0.112 sec  5      0.0498319
55      0.243514      0.02
56 2015-10-30 08:51:52  0.131 sec  6      0.0340885
57      0.19091      0.00666667
58 2015-10-30 08:51:52  0.143 sec  7      0.0241071
59      0.151394      0.00666667
60 2015-10-30 08:51:52  0.153 sec  8      0.017606
61      0.120882      0.00666667
62 2015-10-30 08:51:52  0.165 sec  9      0.0131024
63      0.0975897      0.00666667
64 2015-10-30 08:51:52  0.180 sec  10     0.00976685
65      0.0782481      0.00666667
66
67 Variable Importances:
68 variable      relative_importance      scaled_importance      percentage
69 -----
70 C4      192.761      1      0.774374
71 C3      54.0381      0.280338      0.217086
72 C1      1.35271      0.00701757      0.00543422
73 C2      0.773032      0.00401032      0.00310549

```

5.2.2 Generalized Linear Models (GLM)

Generalized linear models (GLM) are some of the most commonly-used models for many types of data analysis use cases. While some data can be analyzed using linear models, linear models may not be as accurate if the variables are more complex. For example, if the dependent variable has a non-continuous distribution or if the effect of the predictors is not linear, generalized linear models will produce more accurate results than linear models.

Generalized Linear Models (GLM) estimate regression models for outcomes following exponential distributions in general. In addition to the Gaussian (i.e. normal) distribution, these include Poisson, binomial, gamma and Tweedie distributions. Each serves a different purpose and, depending on distribution and link function choice, it can be used either for prediction or classification.

H2O's GLM algorithm fits the generalized linear model with elastic net penalties. The model fitting computation is distributed, extremely fast, and scales extremely well for models with a limited number (\sim low thousands) of predictors with non-zero coefficients.

The algorithm can compute models for a single value of a penalty argument or the full regularization path, similar to glmnet. It can compute Gaussian (linear), logistic, Poisson, and gamma regression models. To generate a generalized linear model for developing linear models for exponential distributions, use `H2OGeneralizedLinearEstimator`. You can apply regularization to the model by adjusting the `lambda` and `alpha` parameters.

```
In [13]: from h2o.estimators.glm import H2OGeneralizedLinearEstimator
In [14]: prostate_data_path = h2o.system_file("prostate.csv")
In [15]: prostate_df = h2o.import_file(path=prostate_data_path)
Parse Progress: [#####] 100%
Imported /Users/hank/PythonEnvs/h2obleeding/bin/../h2o_data/prostate.csv.
  Parsed 380 rows and 9 cols
In [16]: prostate_df["RACE"] = prostate_df["RACE"].asfactor()
In [17]: prostate_df.describe()
Rows:380 Cols:9
Chunk compression summary:
chunk_type  chunk_name                count    count_percentage    size
  size_percentage
-----  -----
CBS          Bits                      1         11.1111             118 B
  1.39381
C1N          1-Byte Integers (w/o NAs)    5         55.5556             2.2 KB
  26.4588
C2           2-Byte Integers              1         11.1111             828 B
  9.7803
CUD          Unique Reals                  1         11.1111             2.1 KB
  25.6556
C8D          64-bit Reals                       1         11.1111             3.0 KB
  36.7116
Frame distribution summary:
                size    number_of_rows    number_of_chunks_per_column
                number_of_chunks
-----  -----
127.0.0.1:54321  8.3 KB    380                1                9
mean            8.3 KB    380                1                9
```

```
29 min                8.3 KB  380                1                9
30 max                8.3 KB  380                1                9
31 stddev             0 B     0                0                0
32 total              8.3 KB  380                1                9
33
34
35
36 In [18]: glm_classifier = H2OGeneralizedLinearEstimator(family="binomial",
37               nfolds=10, alpha=0.5)
38 In [19]: glm_classifier.train(x=["AGE", "RACE", "PSA", "DCAPS"], y="CAPSULE",
39               training_frame=prostate_df)
40 glm Model Build Progress: [#
41               #####] 100%
42 In [20]: glm_classifier
43 Out[20]: Model Details
44 =====
45 H2OGeneralizedLinearEstimator : Generalized Linear Model
46 Model Key: GLM_model_python_1446220160417_6
47
48 GLM Model: summary
49
50     family    link    regularization
51     number_of_predictors_total    number_of_active_predictors
52     number_of_iterations    training_frame
53
54     binomial    logit    Elastic Net (alpha = 0.5, lambda = 3.251E-4 ) 6
55                                     6
56                                     py_3
57
58 ModelMetricsBinomialGLM: glm
59 ** Reported on train data. **
60
61 MSE: 0.202434568594
62 R^2: 0.158344081513
63 LogLoss: 0.59112610879
64 Null degrees of freedom: 379
65 Residual degrees of freedom: 374
66 Null deviance: 512.288840185
67 Residual deviance: 449.25584268
68 AIC: 461.25584268
69 AUC: 0.719098211972
70 Gini: 0.438196423944
71
72 Confusion Matrix (Act/Pred) for max f1 @ threshold = 0.28443600654:
73
74     0      1      Error      Rate
75
76     0      80     147    0.6476    (147.0/227.0)
77     1      19     134    0.1242    (19.0/153.0)
78     Total  99     281    0.4368    (166.0/380.0)
79
80 Maximum Metrics: Maximum metrics at their respective thresholds
81
82 metric                threshold    value    idx
83 -----
84 max f1                0.284436    0.617512    273
85 max f2                0.199001    0.77823    360
```

```
82 max f0point5          0.415159      0.636672  108
83 max accuracy          0.415159      0.705263  108
84 max precision         0.998619      1         0
85 max absolute_MCC      0.415159      0.369123  108
86 max min_per_class_accuracy 0.33266      0.656388  175
87
88 ModelMetricsBinomialGLM: glm
89 ** Reported on cross-validation data. **
90
91 MSE: 0.209974707772
92 R^2: 0.126994679038
93 LogLoss: 0.609520995116
94 Null degrees of freedom: 379
95 Residual degrees of freedom: 373
96 Null deviance: 515.693473211
97 Residual deviance: 463.235956288
98 AIC: 477.235956288
99 AUC: 0.686706400622
100 Gini: 0.373412801244
101
102 Confusion Matrix (Act/Pred) for max f1 @ threshold = 0.326752491231:
103      0      1      Error      Rate
104 ----
105 0      135   92   0.4053   (92.0/227.0)
106 1       48  105   0.3137   (48.0/153.0)
107 Total  183  197   0.3684   (140.0/380.0)
108
109 Maximum Metrics: Maximum metrics at their respective thresholds
110
111 metric                threshold      value      idx
112 -----
113 max f1                 0.326752      0.6        196
114 max f2                 0.234718      0.774359   361
115 max f0point5           0.405529      0.632378   109
116 max accuracy           0.405529      0.702632   109
117 max precision          0.999294      1          0
118 max absolute_MCC       0.405529      0.363357   109
119 max min_per_class_accuracy 0.336043      0.627451   176
120
121 Scoring History:
122 timestamp              duration      iteration      log_likelihood      objective
123 -----
124 2015-10-30 08:53:01    0.000 sec    0              256.482            0.674952
125 2015-10-30 08:53:01    0.004 sec    1              226.784            0.597118
126 2015-10-30 08:53:01    0.005 sec    2              224.716            0.591782
127 2015-10-30 08:53:01    0.005 sec    3              224.629            0.59158
128 2015-10-30 08:53:01    0.005 sec    4              224.628            0.591579
129 2015-10-30 08:53:01    0.006 sec    5              224.628            0.591579
```

5.2.3 K-means

To generate a K-means model for data characterization, use `h2o.kmeans()`. This algorithm does not require a dependent variable.

```

1 In [21]: from h2o.estimators.kmeans import H2OKMeansEstimator
2
3 In [22]: cluster_estimator = H2OKMeansEstimator(k=3)
4
5 In [23]: cluster_estimator.train(x=[0,1,2,3], training_frame=iris_df)
6
7 kmeans Model Build Progress: [#
8     #####] 100%
9
10 In [24]: cluster_estimator
11 Out[24]: Model Details
12 =====
13 H2OKMeansEstimator : K-means
14 Model Key: K-means_model_python_1446220160417_8
15
16 Model Summary:
17   number_of_rows  number_of_clusters  number_of_categorical_columns
18   number_of_iterations  within_cluster_sum_of_squares
19   total_sum_of_squares  between_cluster_sum_of_squares
20
21   -----
22   150                3                0
23   4                405.243                596
24
25 ModelMetricsClustering: kmeans
26 ** Reported on train data. **
27
28 MSE: NaN
29 Total Within Cluster Sum of Square Error: 190.756926265
30 Total Sum of Square Error to Grand Mean: 596.0
31 Between Cluster Sum of Square Error: 405.243073735
32
33 Centroid Statistics:
34   centroid  size  within_cluster_sum_of_squares
35   -----
36   1          96    149.733
37   2          32    17.292
38   3          22    23.7318
39
40 Scoring History:
41   timestamp          duration  iteration  avg_change_of_std_centroids
42   within_cluster_sum_of_squares
43   -----
44   2015-10-30 08:54:39  0.011 sec    0                nan
45   2015-10-30 08:54:39  0.047 sec    1                2.09788
46   2015-10-30 08:54:39  0.049 sec    2                0.00316006
47   2015-10-30 08:54:39  0.050 sec    3                0.000846952
48   190.757

```

5.2.4 Principal Components Analysis (PCA)

To map a set of variables onto a subspace using linear transformations, use `h2o.transforms.decomposition.H2OPCA`. This is the first step in Principal Components Regression.

```

1 In [25]: from h2o.transforms.decomposition import H2OPCA
2
3 In [26]: pca_decomp = H2OPCA(k=2, transform="NONE", pca_method="Power")
4
5 In [27]: pca_decomp.train(x=range(0,4), training_frame=iris_df)
6
7 pca Model Build Progress: [#
8     #####] 100%
9
10 In [28]: pca_decomp
11 Out[28]: Model Details
12 =====
13 H2OPCA : Principal Component Analysis
14 Model Key: PCA_model_python_1446220160417_10
15
16 Importance of components:
17
18 -----
19 Standard deviation      pc1      pc2
20 Proportion of Variance  7.86058  1.45192
21 Cumulative Proportion  0.96543  0.032938
22
23 ModelMetricsPCA: pca
24 ** Reported on train data. **
25
26 MSE: NaN
27
28 In [29]: pred = pca_decomp.predict(iris_df)
29
30 In [30]: pred.head() # Projection results
31 Out[30]:
32      PC1      PC2
33 -----
34 5.9122  2.30344
35 5.57208  1.97383
36 5.44648  2.09653
37 5.43602  1.87168
38 5.87507  2.32935
39 6.47699  2.32553
40 5.51543  2.07156
41 5.85042  2.14948
42 5.15851  1.77643
43 5.64458  1.99191

```


5.3 Grid Search

H2O supports grid search across hyperparameters:

```

1 In [32]: ntrees_opt = [5, 10, 15]
2
3 In [33]: max_depth_opt = [2, 3, 4]
4
5 In [34]: learn_rate_opt = [0.1, 0.2]
6
7 In [35]: hyper_parameters = {"ntrees": ntrees_opt, "max_depth":max_depth_opt,
8                               "learn_rate":learn_rate_opt}
9
10 In [36]: from h2o.grid.grid_search import H2OGridSearch
11
12 In [37]: gs = H2OGridSearch(H2OGradientBoostingEstimator(distribution="
13 multinomial"), hyper_params=hyper_parameters)
14
15 In [38]: gs.train(x=range(0,iris_df.ncol-1), y=iris_df.ncol-1, training_frame
16 =iris_df, nfold=10)
17
18 gbm Grid Build Progress: [#####]
19 100%
20
21 In [39]: print gs.sort_by('logloss', increasing=True)
22
23 Grid Search Results:
24 Model Id                Hyperparameters: ['learn_rate', 'ntrees', '
25 max_depth']          logloss
26 -----
27
28 GBM_model_1446220160417_30  ['0.2, 15, 4']
29                                0.05105
30 GBM_model_1446220160417_27  ['0.2, 15, 3']
31                                0.0551088
32 GBM_model_1446220160417_24  ['0.2, 15, 2']
33                                0.0697714
34 GBM_model_1446220160417_29  ['0.2, 10, 4']
35                                0.103064
36 GBM_model_1446220160417_26  ['0.2, 10, 3']
37                                0.106232
38 GBM_model_1446220160417_23  ['0.2, 10, 2']
39                                0.120161
40 GBM_model_1446220160417_21  ['0.1, 15, 4']
41                                0.170086
42 GBM_model_1446220160417_18  ['0.1, 15, 3']
43                                0.171218
44 GBM_model_1446220160417_15  ['0.1, 15, 2']
45                                0.181186
46 GBM_model_1446220160417_28  ['0.2, 5, 4']
47                                0.275788
48 GBM_model_1446220160417_25  ['0.2, 5, 3']
49                                0.27708
50 GBM_model_1446220160417_22  ['0.2, 5, 2']
51                                0.280413
52 GBM_model_1446220160417_20  ['0.1, 10, 4']
53                                0.28759
54 GBM_model_1446220160417_17  ['0.1, 10, 3']
55                                0.288293
56 GBM_model_1446220160417_14  ['0.1, 10, 2']
57                                0.292993

```

37	GBM_model_1446220160417_16	['0.1, 5, 3']	0.520591
38	GBM_model_1446220160417_19	['0.1, 5, 4']	0.520697
39	GBM_model_1446220160417_13	['0.1, 5, 2']	0.524777

5.4 Integration with scikit-learn

The H2O Python client can be used within scikit-learn pipelines and cross-validation searches. This extends the capabilities of both H2O and scikit-learn.

5.4.1 Pipelines

To create a scikit-learn style pipeline using H2O transformers and estimators:

```

1 In [41]: from h2o.transforms.preprocessing import H2OScaler
2
3 In [42]: from sklearn.pipeline import Pipeline
4
5 In [43]: # Turn off h2o progress bars
6
7 In [44]: h2o.__PROGRESS_BAR__=False
8
9 In [45]: h2o.no_progress()
10
11 In [46]: # build transformation pipeline using sklearn's Pipeline and H2O
           transforms
12
13 In [47]: pipeline = Pipeline([("standardize", H2OScaler()),
14     ....:                      ("pca", H2OPCA(k=2)),
15     ....:                      ("gbm", H2OGradientBoostingEstimator(distribution="
           multinomial"))])
16
17 In [48]: pipeline.fit(iris_df[:4],iris_df[4])
18 Out[48]: Model Details
19 =====
20 H2OPCA : Principal Component Analysis
21 Model Key: PCA_model_python_1446220160417_32
22
23 Importance of components:
24
25 ----- pc1 ----- pc2 -----
26 Standard deviation      3.22082  0.34891
27 Proportion of Variance  0.984534  0.0115538
28 Cumulative Proportion  0.984534  0.996088
29
30
31 ModelMetricsPCA: pca
32 ** Reported on train data. **
33
34 MSE: NaN
35 Model Details
36 =====
37 H2OGradientBoostingEstimator : Gradient Boosting Machine

```

```
38 | Model Key:  GBM_model_python_1446220160417_34
39 |
40 | Model Summary:
41 |   number_of_trees   model_size_in_bytes   min_depth   max_depth
42 |   mean_depth       min_leaves       max_leaves   mean_leaves
43 | -----
44 |   150               27014              1           5           4.84
45 |               2               13           9.99333
46 |
47 | ModelMetricsMultinomial: gbm
48 | ** Reported on train data. **
49 |
50 | MSE: 0.00162796438754
51 | R^2: 0.997558053419
52 | LogLoss: 0.0152718654494
53 |
54 | Confusion Matrix: vertical: actual; across: predicted
55 |
56 | Iris-setosa   Iris-versicolor   Iris-virginica   Error   Rate
57 | -----
58 | 50            0                0              0       0 / 50
59 | 0             50              0              0       0 / 50
60 | 0             0                50             0       0 / 50
61 | 50            50              50             0       0 / 150
62 |
63 | Top-3 Hit Ratios:
64 | k   hit_ratio
65 | ---
66 | 1    1
67 | 2    1
68 | 3    1
69 |
70 | Scoring History:
71 |   timestamp   duration   number_of_trees   training_MSE
72 |   training_logloss   training_classification_error
73 | -----
74 | 2015-10-30 09:00:31  0.007 sec   1.0           0.36363226261
75 |   0.924249463924   0.04
76 | 2015-10-30 09:00:31  0.011 sec   2.0           0.297174376838
77 |   0.788619346614   0.04
78 | 2015-10-30 09:00:31  0.014 sec   3.0           0.242952566898
79 |   0.679995475248   0.04
80 | 2015-10-30 09:00:31  0.017 sec   4.0           0.199051390695
81 |   0.591313594921   0.04
82 | 2015-10-30 09:00:31  0.021 sec   5.0           0.163730865044
83 |   0.517916553872   0.04
84 | ---
85 | 2015-10-30 09:00:31  0.191 sec  46.0           0.00239417625265
86 |   0.0192767794713   0.0
87 | 2015-10-30 09:00:31  0.195 sec  47.0           0.00214164838414
88 |   0.0180720391174   0.0
89 | 2015-10-30 09:00:31  0.198 sec  48.0           0.00197748500569
90 |   0.0171428309311   0.0
91 | 2015-10-30 09:00:31  0.202 sec  49.0           0.00179303578037
92 |   0.0161938228014   0.0
93 | 2015-10-30 09:00:31  0.205 sec  50.0           0.00162796438754
94 |   0.0152718654494   0.0
```

```

84 Variable Importances:
85 variable      relative_importance      scaled_importance      percentage
86 -----
87 PC1           448.958                1                    0.982184
88 PC2           8.1438                 0.0181393           0.0178162
89 Pipeline(steps=[('standardize', <h2o.transforms.preprocessing.H2OScaler
      object at 0x1085cec90>), ('pca', ), ('gbm', )])

```

5.4.2 Randomized Grid Search

To create a scikit-learn style hyperparameter grid search using k-fold cross validation:

```

1 In [57]: from sklearn.grid_search import RandomizedSearchCV
2
3 In [58]: from h2o.cross_validation import H2OKFold
4
5 In [59]: from h2o.model.regression import h2o_r2_score
6
7 In [60]: from sklearn.metrics.scorer import make_scorer
8
9 In [61]: from sklearn.metrics.scorer import make_scorer
10
11 In [62]: params = {"standardize__center":      [True, False],          #
      Parameters to test
12 .....:          "standardize__scale":      [True, False],
13 .....:          "pca__k":                  [2,3],
14 .....:          "gbm__ntrees":              [10,20],
15 .....:          "gbm__max_depth":           [1,2,3],
16 .....:          "gbm__learn_rate":          [0.1,0.2]}
17
18 In [63]: custom_cv = H2OKFold(iris_df, n_folds=5, seed=42)
19
20 In [64]: pipeline = Pipeline([("standardize", H2OScaler()),
21 .....:                        ("pca", H2OPCA(k=2)),
22 .....:                        ("gbm", H2OGradientBoostingEstimator(
      distribution="gaussian"))])
23
24 In [65]: random_search = RandomizedSearchCV(pipeline, params,
25 .....:                                     n_iter=5,
26 .....:                                     scoring=make_scorer(h2o_r2_score)
27 .....:                                     ,
28 .....:                                     cv=custom_cv,
29 .....:                                     random_state=42,
30 .....:                                     n_jobs=1)
31 In [66]: random_search.fit(iris_df[1:], iris_df[0])
32 Out[66]: RandomizedSearchCV(cv=<h2o.cross_validation.H2OKFold instance at 0x108d59200
      >,
33 .....: error_score='raise',
34 .....: estimator=Pipeline(steps=[('standardize', <h2o.transforms.
      preprocessing.H2OScaler object at 0x108d50150>), ('pca', ), ('
      gbm', )]),
35 .....: fit_params={}, iid=True, n_iter=5, n_jobs=1,
36 .....: param_distributions={'pca__k': [2, 3], 'gbm__ntrees': [10, 20], '
      standardize__scale': [True, False], 'gbm__max_depth': [1, 2,
      3], 'standardize__center': [True, False], 'gbm__learn_rate':
      [0.1, 0.2]},

```

```
37         pre_dispatch='2*n_jobs', random_state=42, refit=True,
38         scoring=make_scorer(h2o_r2_score), verbose=0)
39
40 In [67]: print random_search.best_estimator_
41 Model Details
42 =====
43 H2OPCA : Principal Component Analysis
44 Model Key: PCA_model_python_1446220160417_136
45
46 Importance of components:
47 -----
48          pc1          pc2          pc3
49 -----
49 Standard deviation      3.16438      0.180179      0.143787
50 Proportion of Variance  0.994721      0.00322501      0.00205383
51 Cumulative Proportion  0.994721      0.997946      1
52
53
54 ModelMetricsPCA: pca
55 ** Reported on train data. **
56
57 MSE: NaN
58 Model Details
59 =====
60 H2OGradientBoostingEstimator : Gradient Boosting Machine
61 Model Key: GBM_model_python_1446220160417_138
62
63 Model Summary:
64      number_of_trees      model_size_in_bytes      min_depth      max_depth
65      mean_depth      min_leaves      max_leaves      mean_leaves
66  -----
67      20      2743      8      6.35      3      3
68
69 ModelMetricsRegression: gbm
70 ** Reported on train data. **
71
72 MSE: 0.0566740346323
73 R^2: 0.916793146878
74 Mean Residual Deviance: 0.0566740346323
75
76 Scoring History:
77      timestamp      duration      number_of_trees      training_MSE
78      training_deviance
79  -----
79      2015-10-30 09:04:46      0.001 sec      1      0.477453
80      0.477453
80      2015-10-30 09:04:46      0.002 sec      2      0.344635
81      0.344635
81      2015-10-30 09:04:46      0.003 sec      3      0.259176
82      0.259176
82      2015-10-30 09:04:46      0.004 sec      4      0.200125
83      0.200125
83      2015-10-30 09:04:46      0.005 sec      5      0.160051
84      0.160051
84      2015-10-30 09:04:46      0.006 sec      6      0.132315
85      0.132315
85      2015-10-30 09:04:46      0.006 sec      7      0.114554
86      0.114554
```

```
86      2015-10-30 09:04:46 0.007 sec 8          0.100317
      0.100317
87      2015-10-30 09:04:46 0.008 sec 9          0.0890903
      0.0890903
88      2015-10-30 09:04:46 0.009 sec 10         0.0810115
      0.0810115
89      2015-10-30 09:04:46 0.009 sec 11         0.0760616
      0.0760616
90      2015-10-30 09:04:46 0.010 sec 12         0.0725191
      0.0725191
91      2015-10-30 09:04:46 0.011 sec 13         0.0694355
      0.0694355
92      2015-10-30 09:04:46 0.012 sec 14         0.06741
      0.06741
93      2015-10-30 09:04:46 0.012 sec 15         0.0655487
      0.0655487
94      2015-10-30 09:04:46 0.013 sec 16         0.0624041
      0.0624041
95      2015-10-30 09:04:46 0.014 sec 17         0.0615533
      0.0615533
96      2015-10-30 09:04:46 0.015 sec 18         0.058708
      0.058708
97      2015-10-30 09:04:46 0.015 sec 19         0.0579205
      0.0579205
98      2015-10-30 09:04:46 0.016 sec 20         0.056674
      0.056674
99
100 Variable Importances:
101 variable      relative_importance      scaled_importance      percentage
102 -----
103 PC1           237.674                1                0.913474
104 PC3           12.8597                0.0541066        0.0494249
105 PC2           9.65329                0.0406157        0.0371014
106 Pipeline(steps=[('standardize', <h2o.transforms.preprocessing.H2OScaler
      object at 0x104f2a490>), ('pca', ), ('gbm', )])
```

6 References

H2O.ai Team. **H2O website**, 2016. URL <http://h2o.ai>

H2O.ai Team. **H2O documentation**, 2016. URL <http://docs.h2o.ai>

H2O.ai Team. **H2O Python Documentation**, 2015. URL http://h2o-release.s3.amazonaws.com/h2o/latest_stable_Pydoc.html

H2O.ai Team. **H2O GitHub Repository**, 2016. URL <https://github.com/h2oai>

H2O.ai Team. **H2O Datasets**, 2016. URL <http://data.h2o.ai>

H2O.ai Team. **H2O JIRA**, 2016. URL <https://jira.h2o.ai>

H2O.ai Team. **H2Ostream**, 2016. URL <https://groups.google.com/d/forum/h2ostream>